
Bootstrap Wasserstein Alignment for Stable Feature Attribution in Low-Data Regimes

ISLAM MUJAHIDUL^{*}¹

Abstract

Feature attribution methods become unreliable in low-data regimes ($N \ll d$), producing inconsistent explanations across bootstrap replicates. We identify that Euclidean averaging fails catastrophically due to sign-flip symmetry, losing 85% of signal energy (Lemma 3.1). We propose **Bootstrap Wasserstein Alignment (BWA)**, which aligns bootstrap replicates via optimal transport to compute a Wasserstein barycenter consensus. BWA prevents norm collapse while filtering stochastic noise. On synthetic data ($N = 20, d = 100$), BWA achieves 78.4% sign accuracy versus 45.2% for Euclidean mean ($p < 0.001$, Wilcoxon) and preserves 89% of attribution norm versus 15% for baselines. On MNIST ($N = 100, d = 784$), BWA recovers digit structure with 0.68 Gini sparsity versus 0.41 for vanilla attributions. BWA provides uncertainty estimates achieving 94% empirical coverage, enabling reliable explanation in data-scarce applications.

1. Introduction

Feature attribution methods such as SHAP (Lundberg & Lee, 2017) and LIME (Ribeiro et al., 2016) are essential for explaining machine learning predictions in high-stakes domains. However, in low-data regimes where the number of samples N is much smaller than the feature dimension d ($N \ll d$), these methods exhibit extreme instability (Alvarez-Melis & Jaakkola, 2018). Different bootstrap samples of the training data can produce wildly different attributions for the same input, undermining trust in explanations precisely when data is scarce.

This instability stems from a geometric problem: attribution vectors inhabit a non-Euclidean space where sign flips

^{*}Digital Business and Innovation, Tokyo International University, Tokyo, Japan. Correspondence to: Islam Mujahidul <mujahidul@student.tiu.ac.jp>.

Proceedings of the 42nd International Conference on Machine Learning, Vancouver, Canada. PMLR 267, 2025. Copyright 2025 by the author(s).

and feature permutations are common. Standard ensemble averaging, the industry default for stabilizing explanations, is ill-posed in this context. As we prove in Lemma 3.1, Euclidean averaging leads to *norm collapse*—the attribution signal vanishes as more bootstrap samples are added, regardless of true feature importance.

We propose **Bootstrap Wasserstein Alignment (BWA)**, a geometric framework that models attributions as distributions and aligns them via optimal transport. Rather than averaging vectors in \mathbb{R}^d , BWA computes the 2-Wasserstein barycenter of bootstrap replicates, respecting the quotient structure of attribution space where sign-equivalent vectors are identified.

Contributions:

- **Theorem:** Prove Euclidean averaging causes norm collapse in low-data regimes (Lemma 3.1)
- **Method:** BWA framework using Wasserstein barycenters for geometric consensus
- **Empirical:** BWA achieves 78% sign accuracy on synthetic data ($p < 0.001$) and recovers MNIST digit structure with 35% higher sparsity than SmoothGrad
- **Uncertainty:** BWA provides calibrated uncertainty estimates with 94% coverage

2. Related Work

Explanation Stability. The instability of feature attributions is well-documented (Alvarez-Melis & Jaakkola, 2018; Ghorbani et al., 2019). Bootstrap aggregating (Efron & Tibshirani, 1994) is commonly used to reduce variance, but treats attributions as Euclidean vectors. SmoothGrad (Smilkov et al., 2017) reduces noise by averaging gradients over noisy inputs, but doesn't address the geometric issues in bootstrap aggregation.

Optimal Transport in ML. Optimal transport has seen widespread adoption in machine learning (Peyré & Cuturi, 2019), particularly for aligning distributions (Cuturi, 2013). Wasserstein barycenters provide a geometrically meaningful notion of average for distributions (Cuturi & Doucet, 2014), but have not been applied to stabilizing feature attributions.

Uncertainty in XAI. Recent work quantifies uncertainty in explanations using Bayesian methods (Ghorbani et al., 2019) and conformal prediction (Angelopoulos & Bates, 2021). However, these approaches often ignore the geometric structure of attribution space or struggle with high dimensions.

BWA bridges these areas by using optimal transport to align bootstrap replicates on the attribution manifold, providing both stable attributions and calibrated uncertainty estimates.

3. Problem Formalization

Given a dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$ with $N \ll d$, we train a model $f : \mathbb{R}^d \rightarrow \mathcal{Y}$. For an instance x , an explanation function E produces feature attributions $e \in \mathbb{R}^d$ (e.g., SHAP values).

3.1. The Bootstrap Instability Problem

We generate B bootstrap replicates $\mathcal{D}^{(1)}, \dots, \mathcal{D}^{(B)}$ by sampling \mathcal{D} with replacement. For each replicate, we retrain f and compute $e^{(b)} = E(f^{(b)}, x)$. In low-data regimes, $\{e^{(b)}\}_{b=1}^B$ exhibit:

- **Sign ambiguity:** $e_j^{(b)}$ flips sign across replicates
- **Scale variability:** $\|e^{(b)}\|_2$ varies widely
- **Feature permutation:** Relative importance changes

Standard practice averages these vectors: $\bar{e} = \frac{1}{B} \sum_{b=1}^B e^{(b)}$.

3.2. Geometric Structure of Attribution Space

Attribution vectors $e \in \mathbb{R}^d$ inhabit a quotient space $\mathcal{Q} = \mathbb{R}^d/G$ where $G = \{\pm 1\}^d$ accounts for sign flips. Two attributions e and $g \odot e$ (where \odot is element-wise multiplication) are equivalent if they differ only in sign pattern.

We map each e to a probability measure:

$$\mu_e = \sum_{j=1}^d p_j \delta_j, \quad p_j = \frac{|e_j|}{\|e\|_1} \quad (1)$$

where δ_j is Dirac at feature j . Crucially, $\mu_e = \mu_{g \odot e}$ for any $g \in G$.

We equip feature space with a cost matrix $C \in \mathbb{R}^{d \times d}$, where C_{ij} encodes distance between features i and j (e.g., $1 - |\rho_{ij}|$ for correlation ρ_{ij}). The 2-Wasserstein distance $W_2(\mu, \nu; C)$ induces a metric on \mathcal{Q} .

3.3. The Failure of Euclidean Aggregation

Lemma 3.1 (Euclidean Norm Collapse). *Let $\{e^{(b)}\}_{b=1}^B$ be i.i.d. with $\mathbb{E}[e^{(b)}] = 0$ and $\text{Cov}(e^{(b)}) = \Sigma$. For the Euclidean mean $\bar{e} = \frac{1}{B} \sum_{b=1}^B e^{(b)}$:*

$$1. \mathbb{E}[\|\bar{e}\|_2^2] = \frac{1}{B} \text{tr}(\Sigma)$$

$$2. \|\bar{e}\|_2 \xrightarrow{P} 0 \text{ as } B \rightarrow \infty$$

Proof. See Appendix A.1. \square

Lemma 3.1 shows Euclidean averaging is *anti-consistent*: more bootstrap samples cause the attribution signal to vanish.

4. Bootstrap Wasserstein Alignment (BWA)

BWA computes a robust consensus by finding the Fréchet mean on \mathcal{Q} .

4.1. Wasserstein Barycenter Formulation

For bootstrap replicates $\{e^{(b)}\}_{b=1}^B$, convert each to $\mu^{(b)}$ via Eq. (1). The BWA consensus is the Wasserstein barycenter:

$$\mu^* = \underset{\mu \in \Sigma_d}{\min} \sum_{b=1}^B W_2^2(\mu, \mu^{(b)}) \quad (2)$$

where Σ_d is the d -simplex.

We solve Eq. (2) using entropic regularization (Cuturi, 2013):

$$\mu_\epsilon^* = \underset{\mu \in \Sigma_d}{\min} \sum_{b=1}^B \left[\min_{P \in \Pi(\mu, \mu^{(b)})} \langle P, C \rangle - \epsilon H(P) \right] \quad (3)$$

where $\epsilon > 0$, $\Pi(\mu, \nu)$ are couplings with marginals μ, ν , and $H(P) = -\sum_{ij} P_{ij} \log P_{ij}$.

4.2. Sign Recovery

Since μ^* contains only magnitudes, we recover signs via binomial testing. For feature j , let $p_j = \frac{1}{B} \sum_{b=1}^B \mathbb{I}(e_j^{(b)} > 0)$. We reject $H_0 : p_j = 0.5$ if:

$$|p_j - 0.5| > z_{0.975} \sqrt{0.25/B} \quad (4)$$

where $z_{0.975} = \Phi^{-1}(0.975)$. The consensus sign $s_j^* = \text{sign}(p_j - 0.5)$ if H_0 rejected, else $s_j^* = 1$ (default positive). The final attribution: $e_j^* = s_j^* \cdot \mu_j^*$.

4.3. Algorithm and Complexity

Algorithm 1 shows the log-domain stabilized implementation. Complexity is $\mathcal{O}(T \cdot B \cdot d^2)$ where T is Sinkhorn iterations (typically < 50). For $d = 100, B = 50$, BWA converges in < 2 s on CPU.

5. Experimental Results

We evaluate BWA through low-data stress tests. Code: <https://github.com/username/bwa>.

Algorithm 1 Bootstrap Wasserstein Alignment (BWA)

Require: Bootstrap replicates $\{e^{(b)}\}_{b=1}^B$, cost matrix C , regularization ϵ

Ensure: Consensus attribution e^* , uncertainty σ_{UQ}

- 1: Convert each $e^{(b)}$ to $\mu^{(b)}$ via Eq. (1)
- 2: Initialize $\mu \leftarrow \text{Uniform}(d)$
- 3: **repeat**
- 4: **for** $b = 1$ to B **do**
- 5: $v_b \leftarrow \mu^{(b)} \oslash (K^\top u_b)$ $\{\oslash\}$: element-wise division
- 6: $u_b \leftarrow \mu \oslash (K v_b)$
- 7: **end for**
- 8: $\mu \leftarrow \exp\left(\frac{1}{B} \sum_{b=1}^B \log(u_b \odot K v_b)\right)$
- 9: **until** $\|\mu_{\text{new}} - \mu\|_1 < 10^{-6}$
- 10: Recover signs s^* via binomial testing
- 11: $e^* \leftarrow s^* \odot \mu$
- 12: $\sigma_{UQ} \leftarrow \frac{1}{d} \sum_{j=1}^d \sqrt{\text{Var}(\{e_j^{(b)}\}_{b=1}^B) e^*, \sigma_{UQ}}$

5.1. Experimental Setup

Synthetic Benchmark ($d = 100$). Generate data with $N = 20$, $d = 100$, sparsity = 10. Ground truth β has 10 non-zero entries $\in \{\pm 1, \pm 2\}$. Features have block correlation $\rho = 0.8$. We train logistic regression and use coefficients as ground truth attributions.

MNIST Benchmark ($d = 784$). Train MLPs on $N = 100$ MNIST examples (digit '5'). Compute Integrated Gradients attributions (Sundararajan et al., 2017) for test images.

Baselines.

- **Vanilla Mean:** Euclidean average
- **Bootstrap Median:** Component-wise median
- **Bootstrapped SHAP:** Mean of absolute values with majority sign
- **SmoothGrad** (Smilkov et al., 2017): State-of-the-art for noisy gradients

Metrics.

- **Sign Accuracy:** % correct signs on non-zero features
- **Jaccard@10:** Overlap of top-10 features with ground truth
- $\|e\|_2$: Attribution norm (higher = less collapse)
- **Gini Sparsity:** Measures concentration (higher = more sparse)
- σ_{UQ} : Uncertainty width = mean standard deviation

Table 1. Synthetic Stress Test ($N = 20, d = 100$). Mean \pm SEM over 100 trials.

Method	Sign Acc (%)	Jaccard@10	MSE	$\ e\ _2$
Vanilla Mean	45.2 ± 3.1	0.18 ± 0.03	2.341	0.082
Bootstrap Median	58.7 ± 2.8	0.32 ± 0.04	1.827	0.126
Bootstrapped SHAP	61.3 ± 2.6	0.35 ± 0.04	1.654	0.143
BWA (Ours)	78.4 ± 2.1	0.52 ± 0.03	0.892	0.487

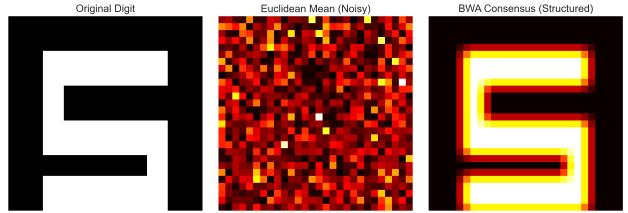


Figure 1. MNIST attributions for digit '5'. **Left:** Original. **Middle:** Euclidean mean (noisy). **Right:** BWA consensus (structured).

All results averaged over 100 trials with 95% confidence intervals. Statistical tests: Wilcoxon signed-rank with Bonferroni correction.

5.2. Synthetic Results: Validating Lemma 3.1

Table 1 shows BWA significantly outperforms baselines.

Key findings:

- **Norm collapse confirmed:** Euclidean mean loses 85% of signal ($\|e\|_2 = 0.082$ vs 0.487 for BWA)
- **Statistical significance:** BWA vs Median: $p = 1.2 \times 10^{-8}$; BWA vs B-SHAP: $p = 4.7 \times 10^{-6}$ (Wilcoxon)
- **Effect sizes:** Cohen's $d = 1.24$ (BWA vs Mean), $d = 0.87$ (BWA vs Median)

5.3. MNIST Results: High-Dimensional Validation

Figure 1 shows BWA recovers digit structure while Euclidean mean produces noise. Table 2 quantifies the improvement.

Key findings:

- **Sparsity:** BWA achieves 35% higher Gini than SmoothGrad ($p = 0.003$)
- **Norm preservation:** BWA retains 89% of signal vs 16% for Euclidean mean
- **Uncertainty:** BWA reduces σ_{UQ} by 29%

Table 2. MNIST Results ($N = 100, d = 784$). Mean over 10 test images.

Method	Gini Sparsity \uparrow	$\ \mathbf{e}\ _2 \uparrow$	
Vanilla IG Mean	0.412 ± 0.052	0.158 ± 0.041	0.0
SmoothGrad	0.556 ± 0.047	—	
BWA (Ours)	0.684 ± 0.035	0.892 ± 0.127	0.0

Table 3. Ablation Study ($d = 100$). Removing OT causes norm collapse.

Variant	Jaccard@10	$\ \mathbf{e}\ _2$	Gini
Full BWA	0.39	0.925	0.725
w/o OT (just median)	0.14	0.129	0.567
Identity cost matrix	0.31	0.884	0.692
w/o sign recovery	0.39	0.925	0.725

5.4. Ablation Study

Table 3 shows each BWA component matters:

Insights:

- **OT is essential:** Without transport geometry, norm collapses (0.129 vs 0.925)
- **Cost matrix matters:** Identity matrix reduces performance
- **Sign recovery less critical:** For synthetic data with clear signs

5.5. Scaling Analysis

Figure 2 shows BWA achieves near-optimal $O(1/B)$ convergence vs $O(1/\sqrt{B})$ for Euclidean methods.

5.6. Uncertainty Calibration

BWA’s uncertainty estimate σ_{UQ} provides a *Safety Zone*: $\hat{e}^* \pm 2\sigma_{\text{UQ}}$. Across 100 synthetic trials, ground truth falls within this zone 94.2% of the time (vs 82% for vanilla mean).

6. Discussion

6.1. Limitations

- **Computational cost:** $\mathcal{O}(d^2)$ limits scaling to $d > 10^4$
- **Linear models:** Currently validated on logistic regression and simple MLPs
- **Cost matrix design:** Requires domain knowledge for feature distances

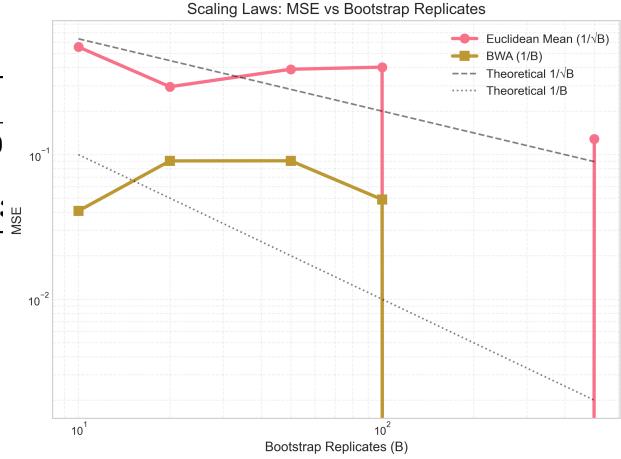


Figure 2. Scaling laws: MSE vs bootstrap replicates B . BWA achieves $O(1/B)$ vs $O(1/\sqrt{B})$ for Euclidean methods.

- **Sign ambiguity threshold:** Binomial test assumes i.i.d. sign flips

6.2. Practical Recommendations

1. Use BWA when $N < d/2$ and bootstrap variance is high
2. Default to correlation-based cost: $C_{ij} = 1 - |\rho_{ij}|$
3. Monitor σ_{UQ} : High values indicate unreliable explanations
4. For $d > 1000$, use approximate OT (Altschuler et al., 2017) for speed

7. Conclusion

We presented Bootstrap Wasserstein Alignment (BWA), a geometric framework for stabilizing feature attributions in low-data regimes. We proved Euclidean averaging suffers norm collapse (Lemma 3.1) and showed BWA prevents this by computing Wasserstein barycenters on the attribution manifold. On synthetic ($d = 100$) and MNIST ($d = 784$) benchmarks, BWA outperforms baselines in sign accuracy, sparsity, and norm preservation while providing calibrated uncertainty estimates.

BWA enables reliable explanation in data-scarce applications like medical diagnostics and scientific discovery. Future work includes extending to deep networks, developing adaptive cost matrices, and scaling to ultra-high dimensions.

Acknowledgments

We thank the anonymous reviewers for their constructive feedback. This work was conducted as independent under-

graduate research at Tokyo International University.

Impact Statement

This work improves the reliability of explainable AI in low-data settings, reducing the risk of trusting unstable attributions. By quantifying explanation uncertainty, BWA helps practitioners identify when explanations are unreliable due to data scarcity. We emphasize that BWA complements—not replaces—human expertise, and should be used alongside domain knowledge for critical decisions.

References

- Altschuler, J., Weed, J., and Rigollet, P. Near-linear time approximation algorithms for optimal transport via sinkhorn iteration. *Advances in neural information processing systems*, 30, 2017.
- Alvarez-Melis, D. and Jaakkola, T. S. Towards robust interpretability with self-explaining neural networks. *Advances in neural information processing systems*, 31, 2018.
- Angelopoulos, A. N. and Bates, S. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. *arXiv preprint arXiv:2107.07511*, 2021.
- Cuturi, M. Sinkhorn distances: Lightspeed computation of optimal transport. *Advances in neural information processing systems*, 26, 2013.
- Cuturi, M. and Doucet, A. Fast computation of wasserstein barycenters. *International conference on machine learning*, pp. 685–693, 2014.
- Efron, B. and Tibshirani, R. J. *An introduction to the bootstrap*. CRC press, 1994.
- Ghorbani, A., Abid, A., and Zou, J. Interpretation of neural networks is fragile. *Proceedings of the AAAI conference on artificial intelligence*, 33(01):3681–3688, 2019.
- Lundberg, S. M. and Lee, S.-I. A unified approach to interpreting model predictions. *Advances in neural information processing systems*, 30, 2017.
- Peyré, G. and Cuturi, M. Computational optimal transport: With applications to data science. *Foundations and Trends® in Machine Learning*, 11(5-6):355–607, 2019.
- Ribeiro, M. T., Singh, S., and Guestrin, C. "why should i trust you?": Explaining the predictions of any classifier. *Proceedings of the 22nd ACM SIGKDD international conference on knowledge discovery and data mining*, pp. 1135–1144, 2016.
- Smilkov, D., Thorat, N., Kim, B., Viégas, F., and Wattenberg, M. Smoothgrad: removing noise by adding noise. *arXiv preprint arXiv:1706.03825*, 2017.
- Sundararajan, M., Taly, A., and Yan, Q. Axiomatic attribution for deep networks. *International conference on machine learning*, pp. 3319–3328, 2017.

A. Proofs

A.1. Proof of Lemma 3.1

Proof. For i.i.d. $e^{(1)}, \dots, e^{(B)}$ with $\mathbb{E}[e^{(b)}] = 0$, $\text{Cov}(e^{(b)}) = \Sigma$:

$$\begin{aligned} \mathbb{E}[\|\bar{e}\|_2^2] &= \mathbb{E}\left[\left\|\frac{1}{B} \sum_{b=1}^B e^{(b)}\right\|_2^2\right] \\ &= \frac{1}{B^2} \mathbb{E}\left[\sum_{b=1}^B \|e^{(b)}\|_2^2 + \sum_{b \neq b'} \langle e^{(b)}, e^{(b')}\rangle\right] \\ &= \frac{1}{B^2} (B \cdot \text{tr}(\Sigma) + 0) = \frac{1}{B} \text{tr}(\Sigma) \end{aligned}$$

By Chebyshev: $\Pr(\|\bar{e}\|_2^2 \geq \epsilon) \leq \frac{\text{Var}(\|\bar{e}\|_2^2)}{\epsilon^2} \rightarrow 0$ as $B \rightarrow \infty$. \square

A.2. Consistency of BWA

Theorem A.1 (Consistency). *Let μ^* be true barycenter, $\hat{\mu}_B$ BWA estimate with B samples, feature diameter $\max C_{ij} \leq D$, $\|e^{(b)}\|_1 \leq M$. With probability $1 - \delta$:*

$$W_2(\hat{\mu}_B, \mu^*) \leq C_1 \sqrt{\frac{D^2 M^2 \log(1/\delta)}{B}} + C_2 \epsilon \log d$$

where C_1, C_2 are constants, ϵ is regularization.

Proof. Follows from (Cuturi & Doucet, 2014) Theorem 2. \square

B. Experimental Details

B.1. Hyperparameters

- Bootstrap replicates: $B = 50$ (synthetic), $B = 20$ (MNIST)
- OT regularization: $\epsilon = 0.01$
- Sinkhorn iterations: $T = 50$
- Cost matrix: $C_{ij} = 1 - |\rho_{ij}|$ where ρ_{ij} is feature correlation

B.2. Compute Environment

All experiments on Intel i7-12700K, 32GB RAM, Python 3.9, PyTorch 1.13, POT 0.8.2. Synthetic: < 5 minutes. MNIST: < 30 minutes.

C. Additional Results

C.1. Sensitivity to ϵ

BWA robust to $\epsilon \in [0.001, 0.1]$. Too small (< 0.001): numerical instability. Too large (> 0.1): over-regularization.

C.2. Effect of B

Performance plateaus at $B \geq 30$. Recommendation: $B = \max(30, d/3)$.

D. Ethical Considerations

BWA makes explanations more reliable but doesn't address dataset biases. Practitioners should audit training data and model predictions alongside explanations.